

Lab 1- Introduction to Socket Programming Part I

1. Sockets

A socket is a conceptual endpoint for a logical network between two processes. i.e., a socket is a 'point' where data can be sent to or received from another process. A socket can send and receive data from a 'remote' process, i.e., a process that runs on a different physical 'host', connected to our host through a computer network.

Recap: TCP and UDP protocols

There are two major protocols used for network traffic in the transport layer in the Internet. They are namely, Transmission Control Protocol (TCP) and User Datagram Protocol (UDP). Both of them reside in the Transport Layer of OSI model. Transport layer is the layer that focuses on process to process communication. i.e., breaking the message into manageable pieces (packets) and multiplexing data between processes.

The difference between the two in the view of the programmer is the reliability of the two. TCP aims to provide an error free and a reliable data transfer while the UDP expects the upper layers to treat the data transfer as inherently unreliable and handle the errors and ordering of packets. To provide the reliability, TCP creates a 'circuit', essentially a path for the packets to travel to the destination. This preserves the ordering of the packets and prevents duplication. Therefore, these are guaranteed when using TCP.

UDP on the other hand does not create such a path. The packets are on their own to figure out a path from the source to destination. This might result in packet drops, duplication of packets and out of order sequencing of packets. Therefore UDP does not provide any guarantee for any of these. Both of these protocols are implemented on top of Internet Protocol (IP). A basic protocol in Internet Layer of the OSI model. This is the layer that focuses on host to host communication. Internet Protocol essentially treats the underlying network as unreliable and does not provide any error checking or sequencing services.

2. Creating a UDP server

A server is the one who waits for a client to connect to. Therefore the client should know the ip address and the port of the server. But the server does not have to know the clients' address. The following steps should be taken to write a UDP server.

Algorithm

1. create a socket with appropriate options for UDP.
2. initialize a `sockaddr_in` struct with appropriate options for our server address.
3. bind the socket with the created `sockaddr_in` struct.
4. send and receive messages from clients that `sockaddr_in` struct is configured to use.

```

/* Sample UDP server */
#include <sys/socket.h>
#include <netinet/in.h>
#include <strings.h>
#include <stdio.h>

int main(int argc, char**argv)
{
    int sockfd,n;
    struct sockaddr_in servaddr, cliaddr;
    socklen_t len; char mesg[1000];
    char* banner = "Hello UDP client! This is UDP server";
    sockfd=socket(AF_INET,SOCK_DGRAM,0);
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr=htonl(INADDR_ANY);
    servaddr.sin_port=htons(32 000);
    bind(sockfd,(struct sockaddr *)&servaddr,sizeof(servaddr));
    len = sizeof(cliaddr); n=recvfrom(sockfd,mesg,1000,0,(struct
sockaddr*)&cliaddr,&len);
    sendto(sockfd,banner,n,0,(struct sockaddr
*)&cliaddr,sizeof(cliaddr));
    mesg[n] = 0;
    printf("Received: %s\n",mesg);
    return 0;
}

```

The socket() call

The socket() function call creates a socket. The return value of this call is an integer called socket descriptor. In UNIX like systems, this is actually a file descriptor and common operations on file descriptors such as read() and write() still work with it. Additionally, specialized functions such as send, sendto(),sendmsg(),recv(),recvfrom() and recvmsg() can also be used with socket descriptors to send and receive data.

To create a socket three arguments are needed. First, a domain. Different domains exist. For example AF_UNIX option is used to communicate in local domain (i.e., inside the local host itself) and AF_INET6 option is used to communicate using IP v6. The option AF_INET specifies that we're using IP v4 to communicate.

The next argument specifies the type of the socket. SOCK_DGRAM specifies that the socket is of type UDP.

The last argument specifies a protocol. This is useful if more than one protocol is supported by the specified type of socket. In this case, there is only one Therefore this number is kept as zero.

Recap: Input output of UNIX systems

Most of the I/O devices are treated as files in UNIX systems. A common interface, `read()` and `write()` is provided to the application programmer by the kernel to deal with these.

The `sockaddr_in` struct

This struct stores IP socket address. An IP socket address contains an IP address of an interface as well as a port number. It contains the following fields:

1. `sin_family`: specifies an address family. Must be `AF_INET`
2. `sin_addr`: stores a struct of `in_port_t` which contains only one field: `s_addr` which stores the IP address in the network byte order
The network and host byte orders and `htonl()` and `htons()` functions
The network and host byte orders and `htonl()` and `htons()` functions
There are two ways that stored in the memory of a computer can be interpreted. They are: the value in the lowest address considered as the most significant byte (named Big Endian) and the value in lowest address considered as the least significant byte (named Little Endian). This is called the 'byte order' of the architecture. The x86 and x86_64 architectures are little endian while the network communication is big endian. Whichever the byte order of the two are, `htonl()` and `htons()` functions converts the byte order of the host machine to the byte order of the network. In the struct, it is the convention to store the port number in the network byte order.
3. `sin_port`: Stores the port to be used in the network byte order.

Bind() call

This gives the socket a 'name' with the given IP address. This basically registers that the socket is using that IP address. Since the server is the initial receiver of the data, this is an essential step. The `bind()` function expects the first argument to be in `struct sockaddr*` type. This is essentially the same as `struct sockaddr_in` in most applications and the casting is done just to get rid of a possible warning.

In the program, you might notice that, `recv_from` has an uninitialized argument `cliaddr` of type `struct sockaddr_in*` as the second argument. The use of `recvfrom` over `read` or `recv` is that, once the data is received, it fills the fields of `cliaddr` with the information of the source the data is received from. This way, the server could figure out the ip address information of the sender and reply back to it. As with `read()`, `recv()` is always blocking.

Creating a UDP client

Algorithm

1. Create a socket
2. Initialize a `sockaddr_in` struct with the information of the server intended to connect to
3. send and receive messages

```

/*Sample UDP client*/

#include<sys/socket.h>
#include<netinet/in.h>
#include <arpa/inet.h>
#include <stdio.h>
#include <string.h>

int main(int argc, char**argv) {
int sockfd,n;
struct sockaddr_in servaddr;
char sendline[] = "Hello UDP server! This is UDP client";
char recvline[1000];
if(argc !=2)
{
printf("usage:%s <IP address>\n",argv[0]);
return -1;
}

sockfd=socket (AF_INET,SOCK_DGRAM,0);
bzero(&servaddr,sizeof(servaddr));
servaddr.sin_family = AF_INET;
servaddr.sin_addr.s_addr=inet_addr(argv[1]);
servaddr.sin_port=htons(32 000);
sendto(sockfd,sendline,strlen(sendline),0,(struct sockaddr
*)&servaddr,sizeof(servaddr));
n=recvfrom(sockfd,recvline,10000,0,NULL,NULL);
recvline[n]=0;
printf("Received: %s\n",recvline);
return 0;
}

```

Unlike the server, the client does not have to bind the socket since the server does not need to know the clients' ip address initially before the client initiate the communication with the server.